

Bridging CPU and GPU in Rust

Niek Aukes Andrei Begu

University of Twente

November 17, 2025

Table of Contents

① Introduction

② Migration Case


③ Compiler Insights

④ Conclusion

Introduction

GPU Model

- GPUs execute thousands of lightweight threads in parallel
- One thread runs the same kernel on different data
- Ideal for data-parallel workloads such as image rendering or simulation
- Heterogeneous computing: best of both worlds!



<https://docs.nvidia.com/cuda/cuda-c-programming-guide/>

Rust

- Modern systems programming language focused on performance and safety
- Enforces memory safety without a garbage collector
- Strong type system and compiler checks
- Built-in tooling: cargo for builds and testing, crates.io for packages
- Previous work on integration with GPUs



Rust (CPU)

```
1 fn vadd(a: &[f32], b: &[f32]) -> Vec<f32> {
2     let mut c = vec![0.0f32; a.len()];
3     for i in 0..a.len() {
4         c[i] = a[i] + b[i];
5     }
6     c
7 }
8
9 fn main() {
10    let a = vec![1.0, 2.0, 3.0, 4.0, ...];
11    let b = vec![10.0, 20.0, 30.0, 40.0, ...];
12    let c = vadd(&a, &b);
13    println!("{:?}", c);
14 }
```

OpenCL I

```
1  use ocl:::{Buffer, ProQue};  
2  
3  const KERN: &str = r#"  
4      __kernel void vadd(__global const float* A,  
5                          __global const float* B,  
6                          __global float* C,  
7                          const int N) {  
8      int i = get_global_id(0);  
9      if (i < N) C[i] = A[i] + B[i];  
10 }"#;  
11  
12 fn main() {  
13     let a = vec![1.0, 2.0, 3.0, 4.0, ...];  
14     let b = vec![10.0, 20.0, 30.0, 40.0, ...];  
15     let n = a.len();  
16     let mut c = vec![0.0f32; n];  
17  
18     let pq = ProQue::builder().src(KERN).dims(n).build().unwrap();
```

OpenCL II

```
19 let a_buf: Buffer<f32> = pq.buffer_builder().copy_host_slice(&a).build().unwrap();
20 let b_buf: Buffer<f32> = pq.buffer_builder().copy_host_slice(&b).build().unwrap();
21 let c_buf: Buffer<f32> = pq.create_buffer().unwrap();
22
23 unsafe {
24     pq.kernel_builder("vadd").arg(&a_buf).arg(&b_buf)
25         .arg(&c_buf).arg(n as i32)
26         .build().unwrap()
27         .enq().unwrap();
28 }
29
30 c_buf.read(&mut c).enq().unwrap();
31 }
```

CUDA C |

```
1  __global__ void vadd(const float* A, const float* B, float* C, int N) {
2      int i = blockIdx.x * blockDim.x + threadIdx.x;
3      if (i < N) C[i] = A[i] + B[i];
4  }
5
6  int main() {
7      float hA[] = {1, 2, 3, 4, ...};
8      float hB[] = {10, 20, 30, 40, ...};
9
10     const size_t N = sizeof(hA)/sizeof(hA[0]);
11     const size_t bytes = N * sizeof(float);
12     float hC[N];
13
14     float *dA, *dB, *dC;
15     cudaMalloc(&dA, bytes);
16     cudaMalloc(&dB, bytes);
17     cudaMalloc(&dC, bytes);
```

CUDA C II

```
18     cudaMemcpy(dA, hA, bytes, cudaMemcpyHostToDevice);
19     cudaMemcpy(dB, hB, bytes, cudaMemcpyHostToDevice);
20
21     vadd<<<128, 4>>>(dA, dB, dC, N);
22     cudaMemcpy(hC, dC, bytes, cudaMemcpyDeviceToHost);
23
24     cudaFree(dA); cudaFree(dB); cudaFree(dC);
25 }
```

Hybrid Rust Compiler

```
1 #[kernel]
2 unsafe fn vadd(a: &[f32], b: &[f32], mut c: Buffer<f32>, n: i32) {
3     let i = gpu::global_tid_x() as i32;
4     if i < n {
5         let sum = a[i as usize] + b[i as usize];
6         c.set(i as usize, sum);
7     }
8 }
9 fn main() {
10     let a = vec![1.0, 2.0, 3.0, 4.0, ...];
11     let b = vec![10.0, 20.0, 30.0, 40.0, ...];
12     let c_buf = Buffer::<f32>::alloc(a.len() as usize).unwrap();
13     let n = a.len();
14
15     vadd.launch(128, 4, &a, &b, c_buf, n).unwrap();
16
17     let result = c_buf.retrieve().unwrap();
18 }
```

Migration Case

Migration Case

- Goal: try-out the compiler on a (small, 300 LOC) real-world case
- Target: existing open-source Mandelbrot renderer in Rust for the CPU
- Per-pixel computation; embarrassingly parallel
- Three stages: migration, testing, maintenance




Figure: Mandelbrot fractal render

Migration (1/4): CPU Version

```
1  for t in 0..threads {
2      thread::spawn(move || {
3          let mut slice = vec![u8::MAX; this_height * image_width];
4          for i in 0..this_height {
5              for j in 0..image_width {
6                  let x = real_start + (j * real_step);
7                  let y = i_start - (i * i_step);
8                  let p = Complex::new(x, y);
9                  if p.is_stable(iterations) {
10                      slice[i * image_width + j] = 0;
11                  }
12              }
13          }
14          tx.send((t, slice)).unwrap();
15      });
16  }
```

Migration (2/4): Extract Function + Annotate

```
1  #[kernel] // annotated as GPU kernel
2  unsafe fn mandelbrot(mut image: Buffer<u8>, offset: usize, ...) {
3      for i in 0..this_height {
4          for j in 0..image_width {
5              let x = real_start + (j * real_step);
6              let y = i_start - (i * i_step);
7              let p = Complex::new(x, y);
8              if p.is_stable(iterations) {
9                  image.set(i * image_width + j, 0);
10             }
11         }
12     }
13 }
```

Migration (3/4): Thread Indexing

```
1 #[kernel]
2 unsafe fn mandelbrot(mut image: Buffer<u8>, offset: usize, ...) {
3     let pos = offset + gpu::global_tid_x();
4     let i = pos / image_width;
5     let j = pos % image_width;
6
7     let x = real_start + (j * real_step);
8     let y = i_start - (i * i_step);
9     let p = Complex::new(x, y);
10    if p.is_stable(iterations) {
11        image.set(i * image_width + j, 0);
12    }
13 }
```

Migration (4/4): Host Launch

```
1 #[kernel]
2 unsafe fn mandelbrot(mut image: Buffer<u8>, offset: usize, ...) {
3     let pos = offset + gpu::global_tid_x();
4     let i = pos / image_width;
5     let j = pos % image_width;
6
7     let x = real_start + (j * real_step);
8     let y = i_start - (i * i_step);
9     let p = Complex::new(x, y);
10    if p.is_stable(iterations) {
11        image.set(i * image_width + j, 0);
12    }
13 }
14
15 mandelbrot.launch(threads, blocks, image, 0, ...).unwrap();
16
```

Migration: The Rest?

**The rest?
The same as before!**

Tooling Project structure

Use of packages Complex type & logic
Image logic Unit tests

Testing

- Existing CPU unit tests run unchanged
- GPU kernels reuse the same tested logic and data types
- One toolchain → regression tests trivial

```
1  #[test]
2  fn test_mandelbrot_gpu_simple_default() {
3      let options = MandelbrotCpu::default();
4      let gpu_image = build_mandelbrot_gpu_simple(&options);
5      let cpu_image = build_mandelbrot_cpu_simple(&options);
6
7      assert_eq!(gpu_image, cpu_image);
8 }
```

Maintenance


- Single-source model reduces long-term maintenance overhead
- CPU and GPU code share the same data structures, types, and logic
- Refactoring affects both sides simultaneously
- All of Rust's toolchain benefits carry over (borrow checks, code smell hints, refactoring tools)
- Can use third party tools in the Rust ecosystem (linters, IDEs, formatters)

Compiler Insights

From Migration to Compilation


- We've seen how Rust can run on GPUs, how does that **actually** work?
- To understand our design, let's look at existing compiler architectures first.

Compiler Architectures: Rust's Perspective




<https://jason-williams.co.uk/posts/a-possible-new-backend-for-rust/>

NVCC: The CUDA Compiler Architecture



<https://hpcgpu.mini.pw.edu.pl/cuda-compilation-toolchain/>

Our Hybrid Compiler Architecture



Advantages and Drawbacks

- Advantages

Advantages and Drawbacks

- Advantages
 - Single pass compiler → no build scripts
 - Single source compilation → no fragmented code

Advantages and Drawbacks

- Advantages
 - Single pass compiler → no build scripts
 - Single source compilation → no fragmented code
 - Safety Guarantees from Rust → including boundary

Advantages and Drawbacks

- Advantages
 - Single pass compiler → no build scripts
 - Single source compilation → no fragmented code
 - Safety Guarantees from Rust → including boundary
- Drawbacks
 - Compiler maintenance and synchronization with `rustc`

Conclusion


Prototype and Future Work

- Current system is a functional prototype

Prototype and Future Work

- Current system is a functional prototype
- Future directions:
 - Update the compiler to a newer version
 - Integration with Rust-GPU
 - Cross-platform support
- Goal: Integrated, cross-platform, safe GPU programming with Rust guarantees

Holk's Rust GPU Compiler (2013)



E. Holk, M. Pathirage, A. Chauhan, A. Lumsdaine and N. D. Matsakis

"GPU Programming in Rust: Implementing High-Level Abstractions in a Systems-Level Language"

Rust's Safety Philosophy

- Memory safety without garbage collection
- Ownership model ensures one writer, many readers
- Borrow checker enforces valid lifetimes and no aliasing
- Type system catches data races and invalid access at compile time

Safety: design goals

- Preserve Rust's memory-safety guarantees inside kernel code
- Make illegal kernel invocations a **compile-time** error
- Make it possible to reason about safety in the **CPU-GPU boundary**
- Restrict which types can be moved to device memory
- Keep unsafe surface minimal and easy to audit

Type-Safe Kernel Invocation

```
1 const vadd: Kernel<u32, (&[f32], &[f32], Buffer<f32>, i32)> = Kernel {  
2     code: &'static [u8],  
3     name: &'static str,  
4     // With Phantom type parameters  
5 }
```

- Compiler encodes argument types in the Kernel type
- Prevents launching kernels with wrong parameters

GPU Memory Management Traits

- **DeepCopy**: asserts a type can be safely `memcpy`'d to the GPU
 - Implemented for plain-old-data types (e.g. `f32`)
 - Guarantees no internal pointers or references

GPU Memory Management Traits

- **DeepCopy**: asserts a type can be safely `memcpy`'d to the GPU
 - Implemented for plain-old-data types (e.g. `f32`)
 - Guarantees no internal pointers or references
- **DSend**: defines how a type is sent to device memory
 - Provides APIs for buffer allocation, transfer, and synchronization
 - Automatically implemented for DeepCopy types
 - Can be manually implemented for complex types (e.g. `&[T]` or `Buffer<T>`)

Pointer Aliasing and Mutable Access

- GPU APIs freely expose raw pointers to shared memory
 - Makes aliasing and data races difficult to reason about
 - No compile-time guarantees of exclusivity

Pointer Aliasing and Mutable Access

- GPU APIs freely expose raw pointers to shared memory
 - Makes aliasing and data races difficult to reason about
 - No compile-time guarantees of exclusivity
- Rust forbids mutable aliasing:
 - One mutable reference (`&mut T`) or many shared (`&T`) references
 - Enforced statically by the borrow checker

Pointer Aliasing and Mutable Access

- GPU APIs freely expose raw pointers to shared memory
 - Makes aliasing and data races difficult to reason about
 - No compile-time guarantees of exclusivity
- Rust forbids mutable aliasing:
 - One mutable reference (`&mut T`) or many shared (`&T`) references
 - Enforced statically by the borrow checker
- Our “unsafe” `Buffer<T>` type provides a controlled escape for GPU memory:
 - Signals that aliasing occurs across host/device boundaries
 - Final safety remains with the programmer

Type Checking through a Library

- Our cuda library governs type-checking and kernel launching
- The library defines:
 - Valid kernel argument types
 - (safe) Memory management rules

Type Checking through a Library

- Our `cuda` library governs type-checking and kernel launching
- The library defines:
 - Valid kernel argument types
 - (safe) Memory management rules
- This enables:
 - Eventually platform-dependent safety semantics (CUDA, Vulkan, ...)
 - Compiler simplicity → Rust's type system does the enforcement
 - Extendability from the community, due to compiler independence

Type-Safe Kernel Invocation

```
1 #[kernel]
2 unsafe fn vadd(a: &[f32], b: &[f32], mut c: Buffer<f32>, n: i32) {
3     let i = gpu::global_tid_x() as i32;
4     if i < n {
5         let sum = a.get(i as usize) + b.get(i as usize);
6         c.set(i as usize, sum);
7     }
8 }
9
10 // Compiles to:
11 const vadd: Kernel<u32, (&[f32], &[f32], Buffer<f32>, i32)> = Kernel {
12     code: &'static [u8],
13     name: &'static str,
14     // With Phantom type parameters
15 }
```